

Interdisciplinary Project (IDP) Master's of Informatics

Chaeun (Joy) Lee, Duc Trung Nguyen

Production Order Release Agent in SAP S/4HANA Public Cloud Edition



Interdisciplinary Project (IDP) Master's of Informatics

Chaeun (Joy) Lee, Duc Trung Nguyen

Production Order Release Agent in SAP S/4HANA Public Cloud Edition

Produktionsauftragsfreigabe-Agent in der SAP S/4HANA Public
Cloud Edition

Examiner


Prof. Dr. Helmut Krcmar

Supervised by

Philipp Landler

Submitted on

30.09.2025



Declaration of Academic Integrity

I confirm that this bachelor's thesis is my own work and I have documented all sources and material used.

I am aware that the thesis in digital form can be examined for the use of unauthorized aid and in order to determine whether the thesis as a whole or parts incorporated in it may be deemed as plagiarism. For the comparison of my work with existing sources I agree that it shall be entered in a database where it shall also remain after examination, to enable comparison with future theses submitted. Further rights of reproduction and usage, however, are not granted here.

This thesis was not previously presented to another examination board and has not been published.

Garching, 30.09.2025

Chaeun (Joy) Lee, Duc Trung

Nguyen

Abstract

This project explores the automation of production order release in SAP S/4HANA Public Cloud Edition through an Artificial Intelligence (AI)-powered agent. Today, supervisors manually check material availability, production capacity, and scheduling across multiple applications, which is slow and error-prone. A major challenge is that Bills of Materials (BOMs) are often incomplete, preventing systems from suggesting alternatives when components are missing.

We extend the Production Order Release Agent with the capability to recommend substitute components. By integrating Joule Functions with OData Application Programming Interfaces (APIs), the agent can detect shortages, propose alternatives from stock or BOM data, and recommend rescheduling if no substitutes exist.

Evaluation indicates that the extended agent reduces manual effort, increases release reliability, and minimizes production downtime. The results highlight how AI agents can improve manufacturing execution in SAP S/4HANA and serve as a practical step toward intelligent, automated production planning.

Contents

Abstract	3
1 Introduction	7
1.1 Context and Motivation	7
1.2 Production Order Domain Terminology	7
General Manufacturing Concepts	7
Production Order Definition and Attributes	7
Production Order Statuses	8
Production Order Process Flow	8
1.3 Problem Statement	8
1.4 Objectives and Research Questions	9
1.5 Structure of the Report	9
2 Related Work	9
2.1 Production Order Release in SAP S/4HANA	9
2.2 SAP Joule	10
2.3 SAP Agent Builder	10
3 Solution Design	11
3.1 Requirements	11
Functional Requirements	11
Non-Functional Requirements	11
3.2 System Architecture	12
Overview	12
3.3 Workflow	12
Primary Workflow	12
Data Flow	13
4 Implementation	13
4.1 Technical Setup	13
4.2 Development Steps	13

	5
4.3 UI Mockups and UX	13
4.4 Challenges During Implementation	13
5 Evaluation	13
5.1 Evaluation Criteria	13
5.2 Test Setup	14
Test Environment	14
Test Dataset	14
Test Scenario Explanations	15
Testing Methodology	15
5.3 Results	16
Overall Performance	16
Component Recommendation Analysis	16
Key Findings	16
5.4 Discussion of Results	17
Strengths of the Agent	17
6 Discussion and Outlook	17
6.1 Interpretation of Results	17
6.2 Challenges and Limitations	17
Environment and Access Constraints	18
Stability and Performance Issues	18
Technical Development Challenges	18
Documentation and Usability Concerns	19
Comparison with Open-Source Alternatives	19
6.3 Future Work	19

List of Tables

1 Production Order Status Definitions	8
2 Test Dataset Overview	14
3 Test Scenario Descriptions	15

List of Figures

Introduction

Context and Motivation

Efficient production order release is a prerequisite for initiating manufacturing execution in SAP S/4HANA. Supervisors must ensure that all required materials and capacities are available before orders are released. This task is traditionally performed through multiple manual checks in different applications, which creates delays, inefficiencies, and risk of human error.

Production Order Domain Terminology

To establish a common understanding of the manufacturing domain, this section defines key terminology related to production orders and factory plant operations.

General Manufacturing Concepts

Work Centers represent specific areas within a plant where manufacturing operations are performed. Each plant typically contains several work centers such as paint, weld, and final assembly stations.

The **Production Supervisor** serves as the mission control for daily production operations, with key responsibilities including:

- Prioritizing order execution
- Solving production problems
- Managing daily production operations

In manufacturing systems, there are typically two types of orders:

- **Planned Orders:** Preliminary production plans
- **Production Orders:** Final instructions to manufacture specific quantities

Production Order Definition and Attributes

A **Production Order** is an instruction to manufacture a specific quantity of a material at a specific time and place. Each production order contains the following key attributes:

- **Material:** What to build

- **Quantity:** How many units to produce
- **Plant + Work Center:** Where the production will occur
- **Start and Finish Date:** When production should begin and end
- **BOM (Bill of Materials):** List of required components
- **Routing:** Which steps and stations to follow
- **Production Version:** Combination of BOM and routing

Production Order Statuses

Production orders progress through various statuses during their lifecycle:

Table 1

Production Order Status Definitions

Status	Description
CRTD	Created but not released
REL	Released for execution
PCNF	Partially confirmed
CNF	Fully confirmed
TECO	Technically completed (no more work expected)
CLSD	Closed (final stage)

Production Order Process Flow

The production order lifecycle follows a structured process flow:

1. **Order Creation:** Convert planned orders to production orders and check material reservation
2. **Release:** Enable material withdrawal for production
3. **Material Staging:** Verify parts availability and issue materials from inventory to shop floor
4. **Execution:** Operators perform manufacturing steps and confirm operations in the system
5. **Good Receipt:** Report finished goods as received and update inventory
6. **Order Settlement & Closure:** Settle costs to cost centers and close the order

Problem Statement

Although SAP S/4HANA Public Cloud already provides ATP checks, capacity validations, and exception handling, these checks rely on accurate and complete BOM data. In practice, production planners often neglect to maintain BOMs properly, leaving gaps in component information. This makes it difficult for existing systems to automatically suggest alternatives when shortages occur, forcing supervisors to manually investigate, cross-check, and reschedule production orders.

Objectives and Research Questions

This IDP develops an AI-powered Production Order Release Agent that goes beyond standard release checks. Specifically, the project extends the agent's capability to:

- Detect missing or unavailable components,
- Propose suitable alternatives from the BOM or current stock using Joule functions,
- Recommend rescheduling if no alternatives are available.

By embedding these features, the agent reduces downtime caused by missing materials and increases the reliability of automated production release.

The project is guided by the following research questions:

1. How can AI agents integrated with SAP Joule Functions and OData APIs support production supervisors in automating order release?
2. How effective is the agent at recommending alternatives when BOM data is incomplete?
3. To what extent does the extended agent reduce manual effort and improve operational efficiency in production order release?

Structure of the Report

The report is organized as follows: Chapter 2 reviews related work in production release automation and AI in manufacturing. Chapter 3 presents the solution design, including system requirements and architecture. Chapter 4 details the implementation in SAP S/4HANA Public Cloud with Joule functions. Chapter 5 evaluates the agent against key performance indicators. Chapter 6 discusses results, interdisciplinary aspects, and limitations. Finally, ?? concludes with key contributions and future work.

Related Work

Production Order Release in SAP S/4HANA

The release of a production order marks the transition from planning to execution in manufacturing. Before release, supervisors must ensure that all required components are available (*availability*

checks) and that machines and labor have sufficient free capacity (*capacity evaluation*). In SAP S/4HANA Public Cloud, these validations are often performed manually across multiple applications, which is time-consuming and error-prone. Although automation can improve efficiency, current systems still assume that the BOM is complete and up to date—an assumption that rarely holds in practice and complicates automated decision-making **sapprodrelease2025**.

SAP Joule

Joule is SAP's generative AI copilot, designed to support business users across different domains of the intelligent enterprise **sapjoule2025**. It provides natural language interfaces, decision recommendations, and integration with enterprise data. In the context of manufacturing, Joule can be extended with custom functions that access transactional data through OData Application Programming Interfaces (APIs). These so-called *Joule Functions* enable developers to implement domain-specific logic, such as retrieving stock levels or proposing alternative materials. SAP provides sample repositories to illustrate how Joule Functions can be created and embedded into workflows, lowering the barrier for extending AI-driven support in S/4HANA.

SAP Agent Builder

The SAP Agent Builder framework allows the creation and orchestration of AI agents within the SAP ecosystem. Agents can encapsulate business logic, interact with core S/4HANA modules, and cooperate with Joule to deliver contextual recommendations. In manufacturing scenarios, Agent Builder offers a structured way to automate repetitive tasks such as production order release checks, while ensuring compliance with enterprise security and governance standards. Previous work has demonstrated the potential of multi-agent systems in industrial settings **wuest2016**, and Agent Builder brings these concepts into the SAP Public Cloud. For this project, it serves as the foundation for developing an AI-powered Production Order Release Agent that can dynamically suggest alternatives and guide supervisors in decision-making.

Solution Design

Requirements

Functional Requirements

The solution design focuses on an AI-enabled production order release system with the following key functional requirements:

- **AI-Enabled Case Management:** The system should leverage artificial intelligence to analyze production order descriptions and compare them with historical data to make intelligent release decisions.
- **Historical Data Analysis:** The system must fetch and analyze production order information from the past month to identify patterns and similarities in component usage and production requirements.
- **Component Availability Checking:** Instead of traditional BOM (Bill of Materials) verification, the system should check past production orders with similar components to assess feasibility.
- **ATP (Available-to-Promise) Validation:** The system must perform ATP checks for missing components to ensure production feasibility before order release.
- **Intelligent Recommendations:** The system should provide alternative solutions based on past events and historical production data.
- **Automated Testing Framework:** Integration with Cucumber framework for comprehensive automated testing of the solution.

Non-Functional Requirements

- **Integration Capability:** Seamless integration with SAP S/4HANA Public Cloud Edition
- **Performance:** Real-time processing of production order release decisions
- **Scalability:** Ability to handle multiple production orders simultaneously
- **Reliability:** High availability and fault tolerance

System Architecture

Overview

The solution architecture is designed around a multi-layered approach that integrates AI capabilities with SAP S/4HANA's existing infrastructure. The core components include:

- **AI Agent Layer:** Central intelligence unit that processes production orders and makes release decisions
- **Joule Integration:** Independent copilot functionality that provides AI-powered assistance
- **OData Services Layer:** Interface for accessing historical production data and SAP system information
- **Testing Framework:** Cucumber-based automated testing infrastructure
- **SAP S/4HANA Integration:** Native integration with the existing SAP system for automated production order release

Workflow

The production order release workflow follows a streamlined process that emphasizes AI-driven decision making:

Primary Workflow

1. **Order Reception:** Production order is received by the system
2. **AI Analysis:** The AI agent analyzes the order description and compares it with historical data
3. **Historical Data Retrieval:** System fetches relevant production order information from the past month
4. **Component Analysis:** Instead of BOM checking, the system analyzes past production orders with similar components
5. **ATP Validation:** Available-to-Promise check for missing components
6. **Recommendation Generation:** AI generates release recommendations based on historical patterns
7. **Alternative Suggestions:** System provides alternative solutions based on past events
8. **Order Release Decision:** Final decision on whether to release the production order

Data Flow

The data flow emphasizes the integration between different system components:

- Production order data flows from SAP S/4HANA to the AI agent
- Historical data is retrieved via OData services from HANA database
- Joule provides additional AI capabilities and tool selection
- Recommendations flow back to the SAP system for final processing

Implementation

Technical Setup

Development Steps

UI Mockups and UX

Challenges During Implementation

Evaluation

Evaluation Criteria

The evaluation of the Production Order Release Agent focuses on three key performance indicators that directly address the research questions outlined in Chapter 1:

- **Accuracy:** The agent's ability to correctly identify missing components and recommend appropriate alternatives
- **Efficiency:** Reduction in manual effort required for production order release decisions
- **Reliability:** Consistency in providing useful recommendations across different scenarios

These criteria are measured against the baseline of manual production order release processes currently used in SAP S/4HANA Public Cloud environments.

Test Setup

Test Environment

The evaluation was conducted using SAP S/4HANA Public Cloud Edition with the following production order management transactions:

- **Create Production Order (CO01):** Used to generate test production orders with various component configurations
- **Change Production Order (CO02):** Applied to modify existing orders and simulate different scenarios
- **Display Production Order (CO03):** Utilized to monitor order status and component availability

These transactions provide comprehensive functionality for production order lifecycle management in SAP S/4HANA, as documented in the official SAP documentation **sapco01**.

Test Dataset

A comprehensive test dataset was created to evaluate the agent's performance across different production order scenarios. The dataset includes 14 distinct production orders with varying component configurations:

Table 2

Test Dataset Overview

Order ID	FrameA	Battery	Wheel	Released	Missing Components
1025282	FrameB	BatteryA	-	N	MATS-BATTERY
1025283	FrameA	BatteryA	-	N	MATS-BATTERY
1025284	FrameA	BatteryA	-	N	MATS-BATTERY
1025285	FrameA	BatteryA	WheelA	N	MATS-BATTERY
1025286	FrameA	-	-	Y	none
1025281	FrameA	BatteryA	-	Y	MATS-BATTERY
1025287	FrameA	BatteryA	WheelA	N	MATS-BATTERY
1025288	FrameB	BatteryA	WheelB	N	MATS-BATTERY
1025289	-	BatteryA	WheelB	N	MATS-FRAME
1025290	-	-	WheelC	N	MATS-FRAME, MATS-BATTERY
1025291	FrameA	BatteryB	WheelA	Y	none
1025292	FrameA	BatteryB	WheelB	N	MATS-WHEEL
1025293	FrameA	BatteryB	WheelC	N	MATS-WHEEL

Test Scenario Explanations

The dataset was designed to test various scenarios and edge cases:

Table 3

Test Scenario Descriptions

Order ID	Test Purpose
1025282	Test agent with different frame type (FrameB) and missing battery
1025283	Baseline case with standard frame and missing battery
1025284	Duplicate scenario to test consistency in recommendations
1025285	Test agent with multiple components including wheel
1025286	Control case: Complete order with no missing components
1025281	Edge case: Released order despite missing battery (manual override)
1025287	Test agent with complete component set but still missing battery
1025288	Test agent with different component combination (FrameB + WheelB)
1025289	Test agent with missing frame component
1025290	Test agent with multiple missing components (frame and battery)
1025291	Control case: Complete order with alternative battery type
1025292	Test agent with missing wheel component
1025293	Test agent with different wheel type and missing component

The dataset was designed to test various scenarios including:

- **Control Cases:** Orders with complete component sets (1025286, 1025291)
- **Missing Component Cases:** Orders with single missing components (1025282-1025285, 1025287-1025288)
- **Multiple Missing Components:** Orders with multiple missing components (1025290)
- **Edge Cases:** Released orders with missing components (1025281)
- **Consistency Testing:** Duplicate scenarios to test recommendation consistency (1025283, 1025284)
- **Variety Testing:** Different component combinations and types

Testing Methodology

The evaluation process involved:

1. **Baseline Measurement:** Recording manual processing time and decision accuracy for each test case
2. **Agent Testing:** Deploying the Production Order Release Agent to process the same test cases

3. **Performance Comparison:** Analyzing differences in processing time, accuracy, and recommendation quality
4. **Error Analysis:** Identifying patterns in cases where the agent provided suboptimal recommendations

Results

Overall Performance

The Production Order Release Agent was tested against the dataset to evaluate its ability to retrieve relevant materials in a more complex dataset across a wider range of time. The evaluation focused on the agent's performance in identifying missing components and recommending appropriate alternatives.

Component Recommendation Analysis

Based on the test results, the agent demonstrated the following performance characteristics:

- **Generally Successful:** The agent was able to retrieve relevant materials in most cases across the test dataset
- **Occasional Redundancy:** In some instances, the agent recommended the same components that were currently missing, indicating areas for improvement in the recommendation logic
- **Complex Dataset Handling:** The agent showed capability in handling the diverse scenarios present in the test dataset

Key Findings

The evaluation revealed that:

1. The agent successfully identified missing components in the majority of test cases
2. Alternative component recommendations were provided for most scenarios
3. Some cases showed redundant recommendations where the agent suggested components that were already identified as missing
4. The agent demonstrated improved performance compared to manual processes in terms of speed and consistency

Discussion of Results

Strengths of the Agent

The evaluation results demonstrate several key strengths of the Production Order Release Agent:

1. **Effective Material Retrieval:** The agent successfully retrieved relevant materials in most test cases, addressing the core challenge of incomplete BOM data
2. **Consistent Performance:** The agent maintained reliable performance across diverse test scenarios
3. **Automated Processing:** The agent reduced manual effort required for production order release decisions

The evaluation demonstrates that the Production Order Release Agent shows promise in addressing the core challenges identified in the problem statement while providing operational benefits in SAP S/4HANA Public Cloud environments. However, there are opportunities for further refinement to improve recommendation accuracy and reduce redundant suggestions.

Discussion and Outlook

Interpretation of Results

How much would it cost to operate the agent in a real-world scenario? <https://www.sap.com/products/data-cloud/hana/pricing.html>

Traces of 9ecf4987-a490-4314-b01f-91f9a7cac630 (Tokens: 47405, Costs: 0.17954 CU)

apply filter on s4 layer to filter for alternative component for Scalability

Challenges and Limitations

The development and deployment of the production order release agent in SAP S/4HANA Public Cloud Edition revealed several significant challenges and limitations that impact both the development process and the practical applicability of the solution.

Environment and Access Constraints

The development environment presented substantial organizational and administrative barriers that significantly reduced development efficiency. Setting up a Business Technology Platform (BTP) subaccount for development required multiple layers of approval and admission processes, creating bottlenecks that limited effective development time. The current setup forces the entire development team to work on a shared subaccount, which severely restricts individual flexibility and creates potential conflicts during concurrent development efforts.

Furthermore, the deployment and testing of new Joule instances proved to be extremely time-consuming due to the requirement for access to both the BTP subaccount and the Canary Cockpit. Creating destinations for agents in the Canary Cockpit requires specific roles that are difficult to obtain, further complicating the development workflow and creating dependencies on administrative approvals that can delay critical development milestones.

Stability and Performance Issues

The Joule platform demonstrated significant stability concerns throughout the development process. The system experienced extended downtime periods, including a nearly two-week outage that severely impacted development progress. Additionally, the response times were consistently slow, particularly during longer "thinking steps" where the agent processes complex reasoning tasks.

A critical limitation emerged regarding the context window capacity, which proved insufficient to handle long purchase order (PO) lists effectively. This constraint directly impacts the practical applicability of the agent in real-world manufacturing scenarios where production orders often involve extensive component lists and detailed specifications.

Technical Development Challenges

The technical development process revealed several significant obstacles. The Handlebar syntax used in Joule proved to be unintuitive and error-prone, increasing development time and the likelihood of implementation errors. OData navigation within the SAP ecosystem was cumbersome, requiring extensive knowledge of SAP-specific data structures and access patterns.

Debugging .yaml scripts became particularly challenging when deployment failures occurred, as the error messages and debugging tools provided limited insight into the root causes of issues.

Additionally, there are restrictions on developing new APIs and limited access to required S/4HANA APIs, which constrains the agent's ability to integrate with all necessary system components.

A fundamental limitation is the absence of a local backend for Joule development. Every request must be rendered on a remote server, making debugging significantly more difficult and increasing development cycle times. This architectural constraint forces developers to rely on remote debugging capabilities, which are often insufficient for complex troubleshooting scenarios.

Documentation and Usability Concerns

The documentation landscape for Joule proved inadequate for effective development. Minimal documentation exists on how Joule actually works internally, forcing developers to rely on trial-and-error approaches and community knowledge. The process of formatting Joule responses correctly for UI5 Cards was particularly unclear, requiring extensive experimentation to achieve proper integration.

Comparison with Open-Source Alternatives

The development experience with Joule contrasts sharply with open-source frameworks such as LangChain. Open-source solutions provide easier local development and testing capabilities, more flexible API and agent integration options, and significantly richer documentation with extensive community support. In contrast, the Joule platform feels closed, restrictive, and slow to iterate on, creating a development experience that hampers rapid prototyping and iterative improvement.

These limitations collectively impact the generalizability of the solution and create dependencies on SAP Cloud infrastructure that may not be suitable for all manufacturing environments. The challenges also raise questions about the long-term maintainability and scalability of the agent-based approach within the current SAP ecosystem constraints.

Future Work